

**SYSC4001 - Assignment 3 Part III**  
**Zaineb Ben Hmida - 101302936**  
**Sohaila Haroun - 101297624**

**1. [0.6 marks]**

Consider the following page reference string in an Operating System with a Demand Paging memory management strategy:

415, 305, 502, 417, 305, 415, 502, 518, 417, 305, 415, 502, 520, 518, 417, 305, 502, 415, 520, 518

**(i) [0.3 marks] Assume we have 3 Frames Allocated. How many page faults will occur with 3 frames allocated to the program using the following page replacement algorithms?**

**a) FIFO (First-In-First-Out)**

415 417 417 **417** 305 305 305 518 518 518 502 502 502 518  
305 **305** 415 415 **415** 502 502 502 417 417 417 415 415 415  
502 502 **502** 518 518 518 520 520 520 305 305 305 520 520

16 page faults

4 page hits **Highlighted**

20 references in total

4/20 = 0.2 therefore 20% hit ratio

**(b) LRU (Least Recently Used) [Student 1: explain the LRU algorithm]**

415 417 417 502 502 305 305 305 520 520 520 305 305 305 520 520  
305 305 305 305 518 518 518 415 415 415 518 518 518 502 502 502 518  
502 502 415 415 415 415 417 417 417 417 502 502 502 417 417 417 415 415

19 page faults

1 hit

20 references in total

1/20 = 0.05 therefore 5% hit ratio

**(c) Optimal**

415 **415** 502 518 518 518 **518** 518 518 **518**  
305 **305** 305 **305** 415 502 520 520 520 **520**  
502 417 417 **417** 417 417 **417** 305 502 415

12 page faults

8 hits **Highlighted**

20 references in total

8/20 = 0.4 therefore 40% hit ratio

**(ii) [0.1 marks] Assume that the case above is repeated with 4 Frames Allocated. Repeat all three algorithms from Part (i) with 4 frames allocated to the program. Show your work for each of the algorithms. Calculate the hit ratio for each of the algorithms**

**a) FIFO (First-In-First-Out)**

415 (Hit 415)	518	518	518 (Hit 518)	518	502
305 (Hit 305)	305 (Hit 305)	415	415	415	415 (Hit 415)
502 (Hit 502)	502	502 (Hit 502)	520	520	520 (Hit 520)
417	417 (Hit 417)	417	417 (Hit 417)	305	305

10 page faults

10 hits

20 references in total

10/20 = 0.5 therefore 50% hit ratio

**(b) LRU (Least Recently Used) [Student 1: explain the LRU algorithm]**

415 (Hit 415)	415	415	305	305	305	518	518	518	518	415	415	415
305 (Hit 305)	305	417	417	417	417	417	520	520	520	520	502	502
502 (Hit 502)	502	502	502	415	415	415	415	417	417	417	417	520
417	518	518	518	518	502	502	502	502	305	305	305	305

17 page faults

3 hits

20 references in total

3/20 = 0.15 therefore 15% hit ratio

**c) Optimal**

415 Hit	415	Hit	502	520	520	520	Hit
305 Hit	305	Hit	305	305	Hit	502	502
502 Hit	518	518	518	518	Hit	518	518
417	417	Hit	417	417	Hit	417	415

6 page faults

11 hits

20 references in total

11/20 = 0.55 therefore 55% hit ratio

**(iii) [0.2 marks] Based on your results, answer the following questions: Which algorithm performs best with 3 frames and why? Which algorithm performs best with 4 frames and why? How do the results change when more frames are allocated? What is the relationship? Why is the Optimal algorithm impractical in real-world operating systems? Compare the performance of FIFO and LRU. When might FIFO be better or worse than LRU?**

**Which algorithm performs best with 3 frames and 4 frames and why?**

The optimal algorithm consistently performs the best when allocating 3 or 4 frames. With a 40% hit ratio, optimal performs best with 3 frames. Compared to FIFO with 20% hit ratio and LRU with a 5% hit ratio, the optimal algorithm performs the best by far. Similar results are seen using 4 frames with a 55% hit ratio optimal algorithm again performs best. Compared to FIFO with a 50% hit ratio and LRU with a 15% hit ratio, the optimal algorithm performs the best among the other algorithms. The optimal algorithm performs the best due to its nature where it always replaces the page that will not be used for the longest amount of time, therefore minimizing the most amount of page faults and resulting in the most hits.

**How do the results change when more frames are allocated? What is the relationship?**

When more frames are allocated, all the algorithms tend to perform better. As can be seen when comparing part i and ii, where FIFO gives a hit ratio of 20% in 3 frames and 50% in 4 frames, LRU gives a hit ratio of 5% in 3 frames and 15% in 4, and finally optimal provides a hit ratio of 40% using 3 frames and 55% using 4. The results show that the hit ratio increases with 4 frames meaning less page faults and more hits and a better overall result.

**Why is the Optimal algorithm impractical in real-world operating systems?**

The optimal algorithm is impractical in real-world operating systems, as you must know the future for it to work. The optimal algorithm works by replacing the page that will not be used for the longest amount of time, meaning the system must know the future references which is not possible in practice. So while in theory it is the best algorithm in practice it is not feasible.

**Compare the performance of FIFO and LRU. When might FIFO be better or worse than LRU?**

Comparing FIFO and LRU, using 3 frames FIFO provides a hit ratio of 20% and LRU one of 5%, and using 4 frames FIFO provides a hit ratio of 50% and LRU one of 15%. These results show that FIFO has a far better performance than LRU in both 3 and 4 frames for this reference string. However, generally LRU performs better in scenarios where recently used pages are highly likely to be used again in the near future which is common in real world systems (principle of locality). FIFO may perform better than LRU in predictable patterns like the ones from this question, where older pages are less likely to be used again.

[Student 1: explain the LRU algorithm]

LRU is the least recently used algorithm, where we replace the page that has not been used for the longest period of time. In the LRU algorithm, when a page must be replaced, the page that has not been used for the longest period of time is replaced by the new page. LRU uses the recent past as an approximation of the future, mimicking the optimal page-replacement algorithm, but looking backward in time rather than forward, which is not feasible in real-world applications.

[Student 2: explain the LFU algorithm]

The Least Frequently Used (LFU) algorithm is a *counting* page replacement algorithm. In counting algorithms, each page has 1-2 bits in the page table that is dedicated to a counter that keeps track of the number of references made to that page. The LFU algorithm then replaces the page that has the smallest count.

**2. [0.3 marks] Consider a system with memory mapping done on a page basis. Assume that the necessary page table is always in main memory. A single main memory access takes 120 nanoseconds (ns).**

**(a) [0.1 marks] How long does a paged memory reference take in this system without a TLB? Explain your answer.**

Memory access time = page table + data

Memory access time = 120ns + 120ns = = 240ns

In a system with memory mapping done on a page basis every memory access first reads the page table (always in main memory for this example) to find the address in memory and then access the data in that address in memory. Since each memory access takes 120 ns, the total time for a paged memory reference without a TLB is 240ns.

**(b) [0.1 marks] If we add a Translation Lookaside Buffer (TLB) that imposes an overhead ( $t$ ) of 20 ns on a hit or a miss. If we assume a TLB hit ratio of 95%, what is the effective memory access time (EMAT)? Explain your answer.**

$$\text{EMAT} = h(t + m) + (1 - h)(t + m + m)$$

$$\text{EMAT} = 0.95(20 + 120) + (1 - 0.95)(20 + 120 + 120)$$

$$\text{EMAT} = (0.95 \times 140) + (0.05 \times 260)$$

$$\text{EMAT} = 146\text{ns}$$

A Translation Lookaside Buffer (TLB) caches recent page table entries. Adding a TLB means that on a hit, we no longer require reading the page table from main memory, but it imposes an overhead of 20ns per lookup. With a 95% hit ratio, most of the time, the page table does not need to be read from main memory, but on a miss, the memory system must read the page

table from main memory and the data, plus the 20ns overhead per lookup. The effective memory access (EMAT) is therefore 146ns.

**(c) [0.1 marks] Why does adding an extra layer, the TLB, generally improve performance?**

**Are there situations where the performance may be worse with a TLB than without one? Explain all cases.**

Adding an extra layer, the TLB, generally improves performance because it allows some memory access to avoid the main memory page lookup step, which takes time as it is an extra main memory access per main memory access, and this reduces the total amount of memory operations required overall (if the hit ratio is good). This is especially beneficial when the same page is accessed repeatedly. Situations where the performance may be worse with a TLB than without one are when the TLB hit rate is very low and the overhead imposed by the TLB is high. Consider the case where the page's entry is not in the TLB, this memory access will both require 2 main memory accesses and has already been affected by the TLB overhead ( $t + 2m$ ). This results in slower memory accesses when using both the TLB and the normal memory access, resulting in a larger effective access time than a system without a TLB.

**3. [0.3 marks] Consider a system with a paged logical address space composed of 128 pages of 4Kbytes each, mapped into a 512 Kbytes physical memory space. Answer the following questions and justify your answers.**

**(a) [0.1 marks] What is the format and size (in bits) of the processor's logical address?**

Logical address space = 128 pages

Pages = 4 Kbytes

Pages =  $4 \times 1024 = 4096$  bytes

Number of bits per page number =  $\log_2(\text{logical address space})$

Number of bits per page number =  $\log_2(128)$

Number of bits per page number = 7 bits

Number of bits per offset =  $\log_2(\text{pages})$

Number of bits per offset =  $\log_2(4096)$

Number of bits per offset = 12 bits

Processor's logical address size =  $7 + 12$

Processor's logical address size = 19 bits

Processor's logical address format:

[ 7-bit page number | 12-bit offset ]

**(b) [0.1 marks] What is the required length (number of entries) and width (size of each entry in bits, disregarding control bits) of the page table?**

Logical address space = 128 pages  
Pages = 4 Kbytes  
Pages =  $4 \times 1024 = 4096$  bytes  
Physical memory space = 512 Kbytes  
Physical memory space =  $512 \times 1024 = 524,288$  bytes

Number of frames = Physical memory space ÷ Pages  
Number of frames =  $524,288 \div 4096$  or (512KB/4KB)  
Number of frames = 128 frames

Number of bits for the frame =  $\log_2(\text{Logical address space})$   
Number of bits for the frame =  $\log_2(128)$   
Number of bits for the frame = 7 bits

Therefore, the required length (number of entries) of the page table is 128 pages, and the required width (size of each entry in bits, disregarding control bits) of the page table is 7 bits.

**(c) [0.1 marks] What is the effect on the page table width if now the physical memory space is reduced by half (from 512 Kbytes to 256 Kbytes)? Assume that the number of page entries and page size remain the same.**

If the physical memory is halved, the page table width will be 1 bit smaller. (this is clear just from understanding how bits work but here are some calculations to prove it:)

Number of frames = Physical memory space ÷ Pages  
Number of frames = 256KB/4KB  
Number of frames = 64 frames

Number of bits for the frame =  $\log_2(\text{Logical address space})$   
Number of bits for the frame =  $\log_2(64)$   
Number of bits for the frame = 6 bits

[Student 1: explain what the physical memory space is]

Physical memory space is the actual hardware memory, RAM, of the computer. Physical memory space refers to the limited, real memory location that the CPU can directly access. The physical memory is divided by the operating system into frames, or fixed-size memory blocks, used to work with virtual memory. Data in physical memory is volatile and is lost when the CPU powers off.

[Student 2: explain what the logical memory space is]

The logical memory space is the set of addresses that a program *thinks* it has while it is running.

These addresses (logical/virtual) are handled by the CPU when the program executes instructions.

Logical memory is independent of where the program is actually stored in physical RAM.

In a virtual memory system, this exists because of virtual memory, which separates the program's view of memory (logical addresses) from the actual hardware memory (physical addresses). Because of this separation:

- A program does not need to have all its code in RAM at once—only the parts being executed.
- The logical address space can be much larger than the physical RAM, since virtual memory can use disk space to extend it.
- Multiple processes can each have their own logical address space, making it appear as if each process has “its own private memory”.

Quick summary (mostly for myself):

The compiler generates relative addresses (like using offsets).

When a program runs, the CPU generates/uses logical addresses (also known as virtual addresses).

The MMU translates logical addresses → physical addresses (for ex: using page tables).

**4. [0.1 marks] Explain, in detail, the sequence of operations and file system data structure accesses that occur when a process executes the lseek(fd, offset, SEEK\_END) system call.**

**Consider a system using a hierarchical directory structure and assume the file described by the file descriptor (fd) is not currently open by any other process.**

When a process executes the lseek(fd, offset, SEEK\_END) system call:

It creates a “software interrupt” which causes a switch to kernel mode and a context switch. The provided interrupt vector is then used to get the first address of the corresponding ISR from the vector table, which will be the first address of the lseek routine.

The operating system then uses the file descriptor to look up the corresponding entry in the process's file descriptor table, which points to the system-wide open file table entry for that file. From there, the OS accesses the file's inode, which contains the current file size which is needed because SEEK\_END means the new position is calculated relative to the end of the file. If the inode is not already cached, it must be read from disk. The kernel then computes the new offset as file size + offset and updates the stored file position in the open file table entry. Importantly, the directory structure is not accessed during this call because the file is already open, and no data blocks are touched until the process actually reads or writes. Finally, the new position is returned to the process.

## 5. File System Organization

a) [0.1 marks] (from Silberschatz) Consider a file system that uses inodes to represent files. Disk blocks are 8Kb in size, and a pointer to a disk block requires 4 bytes. This file system has 12 direct disk blocks, as well as single, double, and triple indirect disk blocks. What is the maximum size of a file that can be stored in this file system?

Direct blocks:

$$12 * 8\text{KB} = 96\text{KB}$$

Single indirect:

$$\begin{aligned} 1 \text{ block}, 2^{11} &= 2048 \text{ pointers} \rightarrow 2048 * 8 \text{ KB} \\ &= 16,384 \text{ KB} = 16 \text{ MB} \end{aligned}$$

Double indirect:

$$\begin{aligned} 2048 \text{ (first level)} * 2048 \text{ (second level)} * 8 \text{ KB} \\ &= 4,194,304 \text{ blocks} * 8 \text{ KB} \\ &\approx 32 \text{ GB} \end{aligned}$$

Triple indirect:

$$\begin{aligned} 2048^3 * 8 \text{ KB} \\ &= 8,589,934,592 \text{ blocks} * 8 \text{ KB} \\ &= \sim 68,719,476,736 \text{ KB} \approx 64 \text{ TB} (* \text{ this doesn't feel right, but I tried following with the textbook and got lost}) \end{aligned}$$

Direct:  $12 \times 8192 = 98,304 \text{ bytes} (\approx 96 \text{ KB})$

Single indirect:  $2048 \times 8192 = 16,777,216 \text{ bytes} (\approx 16 \text{ MB})$

Double indirect:  $2048^2 \times 8192 = 34,359,738,368 \text{ bytes} (\approx 32 \text{ GB})$

Triple indirect:  $2048^3 \times 8192 = 70,368,744,177,664 \text{ bytes} (\approx 64 \text{ TB})$

Total maximum file size = 96 KB + 16 MB + 32 GB + 64 TB

Total  $\approx 64 \text{ TB}$  (dominated by triple indirect)

Maximum file size  $\approx 64 \text{ TB}$ .

Or is it:

direct =  $12 * 8\text{KB} = 48 \text{ KB} = 384000$

indirect =  $12 * 8\text{KB} * 8\text{KB}$

Double indirect =  $12 * 8\text{KB} * 8\text{KB} * 8\text{KB}$

Triple =  $12 * 8\text{KB} * 8\text{KB} * 8\text{KB} * 8\text{KB}$

**b) [0.1 marks] Explain what you can do in case (a) if you need to store a file that is larger than the maximum size computed. Give an example showing how you can define a larger file, and what the size of that file would be.**

If you need a larger file you can (A) split the logical file across multiple files and manage them at the application level or with a volume layer, (B) reformat the filesystem with a larger block size (or use a filesystem with wider pointers/offsets), or (C) use a distributed/cluster filesystem or logical volume that presents a larger address space.

If the on-disk inode structure can address more space, the kernel's in-memory file-offset width (e.g., 32-bit signed) may still limit usable file size to  $2^{31}-1$  bytes unless the OS/filesystem uses 64-bit offsets. (according to textbook)